Inspectable Narrative Graphs
Local Anonymous Companion

Retained evidence, requests, outputs and assessments

1. Reading the networks

This supplement adds connected graph views, a temporal explanation and a worked error analysis to the main paper. It is not a second experiment. All outputs, reference answers and assessments are retained from the two reported batches. Codex authored the reference annotations and semantic judgments. These are not independent human review.

Two paths. Pre-extract/select (A) selects existing records from a model extraction made without the question. Contextual extraction (B) independently receives the text and question. A and B graphs are separate outputs, not stages of the same graph.

How to read an edge. Arrows retain subject-to-object direction. Labels show the actual relation (underscores displayed as spaces), citations and any stated bounds or holder/attitude. `loc` abbreviates `located_in`. Identical endpoint strings share a node within a panel. Different strings are never merged for appearance. Parallel assertions remain separate edges. Unprinted qualifiers are null, not timeless truth. A holder badge is a qualification, not an extra model-generated relationship.

Assessment and relevance. + denotes supported meaning, ∼ partial meaning, x unsupported and ? unresolved. I denotes supported but irrelevant material. Colour is redundant with labels and line styles. Node fills in Orchard distinguish people, objects and places as display annotations, not inferred ontology types. The broad main-paper extraction is deliberately ungraded rather than presented as a reference graph.

Compact index of retained examples

Example	What to inspect	Complete record location
Harbor	Source intervals versus query window (p. 4)	Compact calls 1, 4–6
Orchard	Connected ownership slice, extra locations and attribution (pp. 2–3)	Compact calls 2, 7–9
Museum	Fresh synthetic comparison, including relevance errors	Compact calls 3, 10–12
Fable	Parallel actions, passive wording and conditional scope (pp. 5–6)	Prose calls 1, 4–5
Alice	Wrong reading participant and literal thought object (p. 6 and main paper)	Prose calls 2, 6–7
Holmes	Syntax failure and unresolved interpretation (pp. 6–7)	Prose calls 3, 8–9

Call numbers are local to their batch. Three pre-extractions precede contextual requests in each batch. A pre-extraction reused for several questions is still one model call. The main paper contains the complete aggregate tables. This supplement retains selected explanations instead of repeating every prompt and record. Page 7 gives exact file and field paths for the complete data.

The main Orchard figure is an illustrative selection from an existing evaluation, not a prospectively sampled new success. The views below include extra irrelevant predictions, partial meanings and failures. No record was repaired for display, and no score was recomputed.

2. Orchard: correct relationships can still be irrelevant

Exact executed ownership/carrying question: Who owns or carries what? Return all and only directly narrated ownership and carrying relationships, retaining every stated interval.

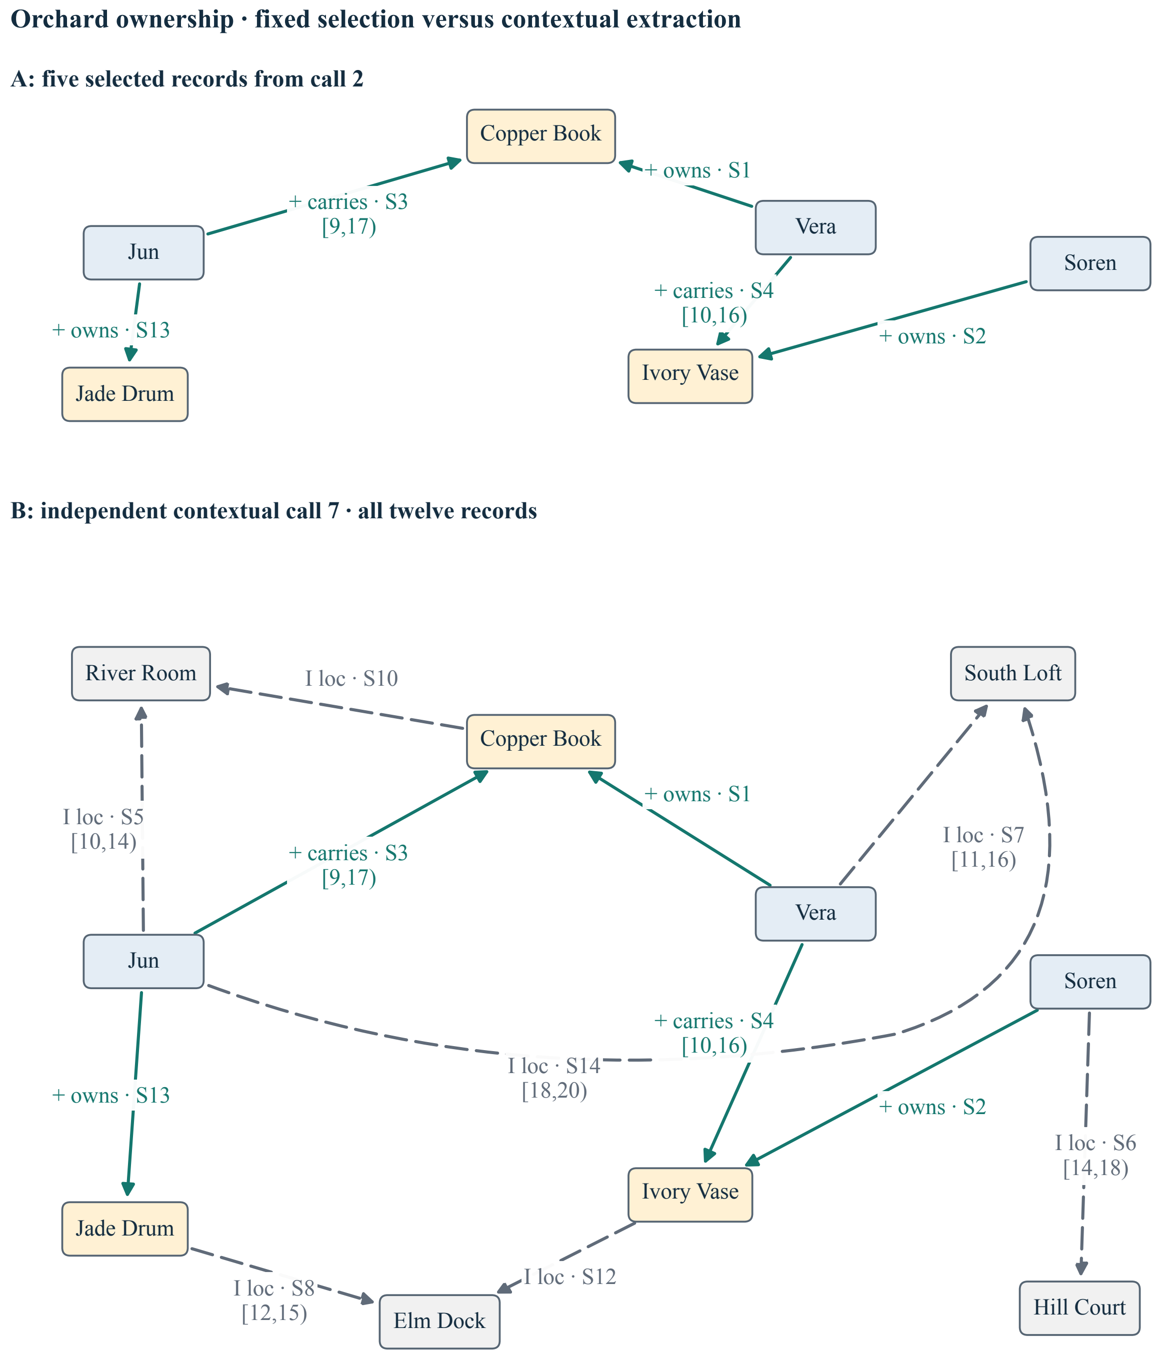

+ supported and relevant. I supported but irrelevant (dashed). Colours assess records, not selection. Node fills are display categories. loc = located in. All holder/attitude values null. Unprinted bounds null.

A: 5/5 matched, 5 targets. P/R/F1 = 1.000/1.000/1.000. **B:** 5/12 matched, 5 targets. P/R/F1 = 0.417/1.000/0.588.

Both views preserve all five requested facts. The contextual output also returns seven locations. These are supported by the supplied story but outside this question. They remain visible and in the precision denominator. No belief records are present in this contextual ownership output. The panels share horizontal anchors and node categories, not an inferred identity mapping across different strings.

3. Orchard: attribution belongs to the assertion

Exact executed question: Who explicitly believes what, and what reality is separately narrated about those same objects? Return each underlying believed relationship with its attribution and each matching narrated reality without attribution. Omit unrelated facts.

Orchard · believed and narrated locations

S9 Soren believes that Copper Book is in South Loft.

S10 In reality, Copper Book is in River Room.

S11 Jun believes that Ivory Vase is in Hill Court.

S12 In reality, Ivory Vase is in Elm Dock.

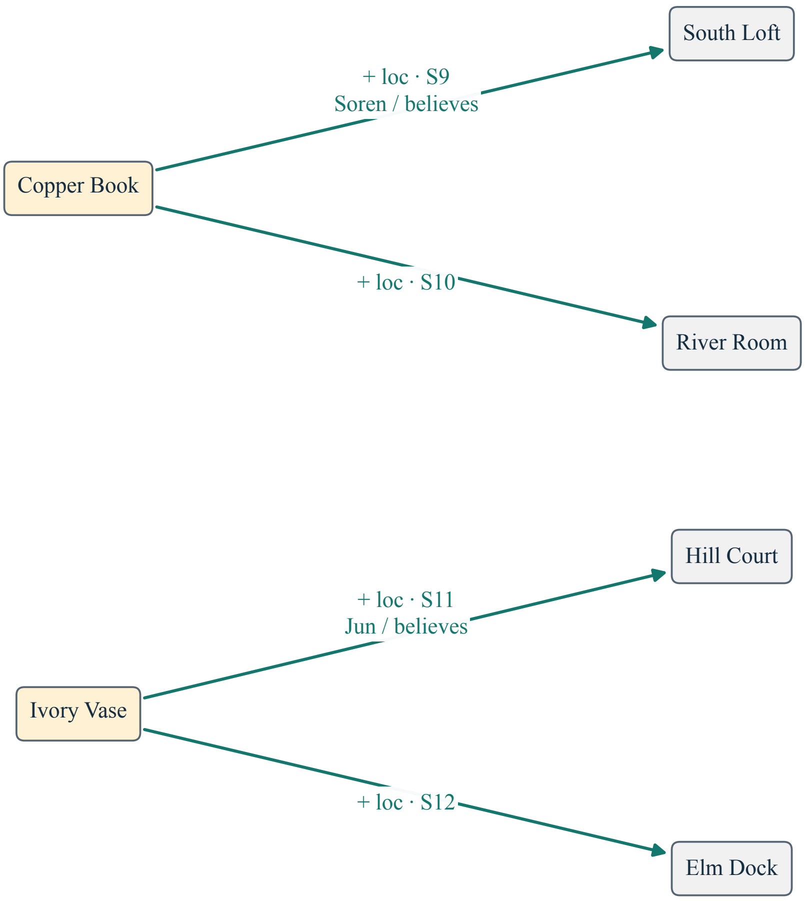

A: fixed selection from call 2. All four records on six endpoint nodes. Holder names are qualifications, not additional nodes or edges. All validity bounds null. + supported under the frozen evaluation. loc = located in.

Fixed-selection score: 4/4 matched, 4 targets. P/R/F1 = 1.000/1.000/1.000.

Each object connects to two places under different epistemic scopes. Copper Book's South Loft location is attributed to Soren's belief, while River Room is narrated reality. Ivory Vase similarly separates Jun's belief from narration. Holder names remain edge metadata even though the names also occur elsewhere in the broader extraction. This selected graph has two disconnected components. No connector has been invented to make it appear more cohesive.

4. Harbor: preserve source duration, select by overlap

Exact executed question: Which directly narrated location relationships have an explicitly stated interval overlapping day 2 up to but not including day 4? Return all of them with their full evidence-stated intervals, not clipped to the question window. Exclude undated locations and beliefs.

Harbor · full intervals, not clipped to the question

S5 Lena is in North Hall from day 1 up to but not including day 5.

S6 Omar is in East Garden from day 0 up to but not including day 3.

S7 Priya is in Stone Shed from day 3 up to but not including day 6.

S8 Silver Flute is in West Pier from day 2 up to but not including day 4.

Request: narrated locations overlapping [2,4), with full source intervals.

A and B: four identical location records, generated independently

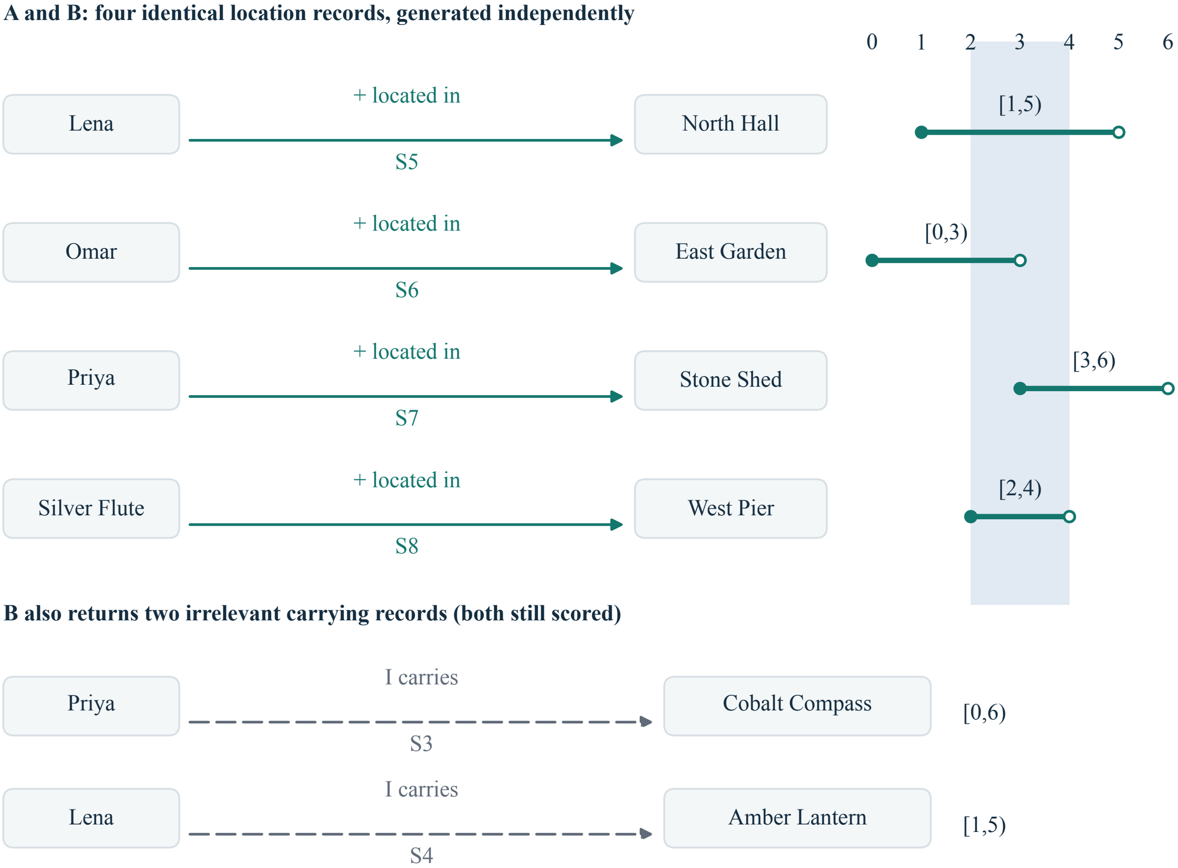

+ supported I irrelevant; all holder/attitude values null. Shading = query window; filled start / open end = source validity.

The blue band is the requested window $[2,4)$, not rewritten validity. Filled starts and open ends mark the full generated intervals. All four selected A location records coincide with four B records. The two carrying records below are additional B predictions, not part of the location answer. All six B records remain scored. No relationship is given extra temporal precision because it was observed during the question window.

Source display: S5–S8 of the fourteen supplied sentences. Complete evidence and both outputs remain in the compact records. The main paper reports their unchanged strict scores in Table 1.

5. Fable: parallel actions and separate endpoint identities

Fable actions · full networks with literal endpoint identities

A: fixed selection · all eight records

B: independent contextual extraction · all five records

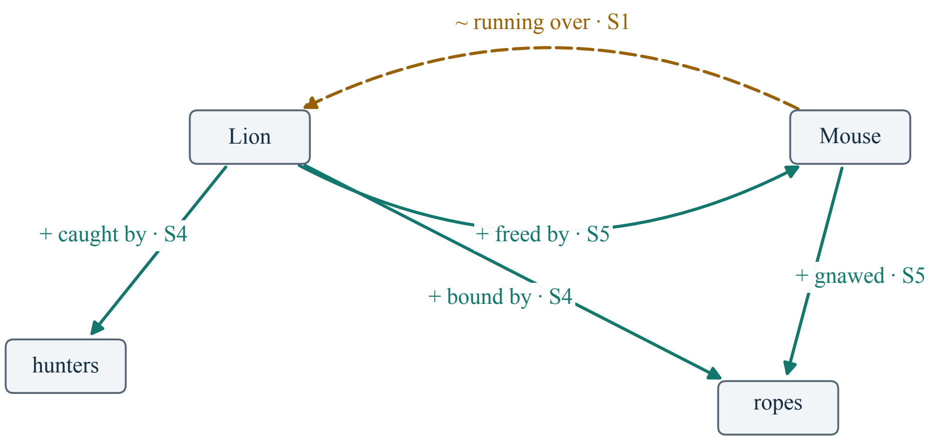

+ supported meaning. ~ partial. ? unresolved. Exact endpoint strings remain separate: A has “strong ropes” and “the rope”, whereas B uses “ropes”. All attribution and bounds null except A’s missing valid_until field on spare (literal extra key “valid until”).

Same executed question as main Figure 2: Which physical running/waking, capture, binding, rope-gnawing and release relationships involve the Lion, Mouse, hunters and ropes? Exclude dialogue, intentions, laughter and the moral.

A: P/R/F1 = 0.625/0.625/0.625. **B:** P/R/F1 = 0.400/0.250/0.308.

All eight A and five B action records appear. A’s five nodes preserve its two different rope strings. B’s four nodes use a single literal ropes endpoint. This is not a corrected alias alignment. The complete source is the 133-word fable retained in the prose dataset. A’s ambiguous awakening and partial conditional plea remain in the graph.

6. Agreement, preserved meaning and substantive errors

A worked example from the fable

Exact supplied S4: It happened shortly after this that the Lion was caught by some hunters, who bound him by strong ropes to the ground.

Actual contextual record: Lion | caught_by | hunters | S4. Its bounds, holder and attitude are all null. The retained assessment is:

> S4 supports this passive capture relation exactly in meaning. The frozen strict matcher lacks the inverse caught_by alternative; its strict failure is not a falsehood judgment.

The passive relationship preserves the narrated participants and meaning. The frozen strict matcher does not accept that inverse wording. This is a strict mismatch without a semantic falsehood verdict. It is not changed into a matching record for these figures.

Exact supplied S1: A LION was awakened from sleep by a Mouse running over his face.

Actual contextual record: Mouse | running_over | Lion | S1. The retained assessment is:

> S1 says running over the Lion's face. Running over Lion retains the actor and broad action but loses the explicit body-part endpoint. This receives partial support, not full coverage of the frozen face-specific target.

The broad action survives, but the face-specific endpoint does not. Partial support therefore does not count as complete target coverage. Neither this case nor the passive example is newly scored.

Compact failure and uncertainty guide

Issue	Retained example	Interpretation
Wrong participant	Alice contextual actions, record 4: Alice reading the book	Her sister reads. Alice peeps into the book. This is unsupported participant attribution, not merely a synonym mismatch.
Belief representation	Alice contextual claims, record 2: literal thought object	Meaning is supported, while the requested decomposed structure is absent. The daisy-chain record remains unresolved.
Irrelevant additions	Orchard contextual possession, seven locations	Supported relationships outside the question. Counted as precision errors, not silently removed.
Conditional scope	Fable A action record 3: spare	The Mouse's conditional plea is not an unconditional release. The misspelled temporal key also remains.
Syntax failure	Holmes pre-extraction and contextual actions	Extra closing brace. No parsed graph scores and no invented empty network.

Pending author judgments remain pending, including Holmes's vocative and speaker attribution. Unresolved wording is not assumed true or false.

7. Reproducibility without repeated data dumps

Complete records remain available

Paths below are relative to the conference package unless prefixed by *repository*. They contain the full information that the former 36-page companion printed repeatedly. The data are unchanged.

File	What to read
`data/requests.json`	Exact transmitted payloads, including system and user messages. Select `compact` or `prose`, then the string call ID. Request hashes are validated.
`data/retained_outputs.json`	Complete evidence under `stories`. Exact raw response strings and parsed records under `calls`. Full frozen reference alternatives under `references`. Actual selected/contextual records and evaluations under `results`.
`data/evaluation_rules.json`	Frozen normalization, scope and matching rules for each batch.
`manifest.json`	Exact displayed facts, assessment labels, source call/record mappings, endpoint identities and layout coordinates. Display subsets do not change denominators.
`generated/reproducibility_index.json`	Hashes and field-level paths connecting this reader guide to complete retained evidence and historical companion version.
Repository `reports/tables/real_text_proof_of_concept.json`	Complete Codex semantic judgments under `results[].semantic_assessments`, plus pre-extraction assessments and syntax diagnoses.
Repository `reports/tables/compact_story_v2_results.json`	Complete synthetic results, source-support evaluations, selection traces and all historical syntax-recovery categories.

The published prose is preserved verbatim with section, edition, retrieval and hash metadata. Project Gutenberg notices remain in `data/pg21_notices.txt`, `data/pg11_notices.txt` and `data/pg1661_notices.txt`. Aesop uses Townsend's translation of *The Lion and the Mouse*; Alice is from Chapter I, and Holmes from *The Red-Headed League*. Neither familiar context nor later plot information is used to resolve a model record.

A syntax failure is not an empty graph

Holmes contextual actions, retained call 8. Display excerpt: the final 200 characters, unchanged:


```
l,"valid_until":null,"holder":null,"attitude":null},{"subject":"Holmes","relation":"
 ↪  said","object":"Watson","evidence_ids":["S7"],"valid_from":null,"valid_until":nu
 ↪  ll,"holder":null,"attitude":null}]}}
```


Original parser diagnosis: Extra data: line 1 column 1683 (char 1682)

Only the display is excerpted. The complete raw response remains in `data/prose_call_8_raw.txt` and the JSON record. The narrow permitted trailing-comma recovery does not remove an extra closing brace. No graph is reconstructed here.

Provenance and access

The previous 36-page companion remains in version history. There is one current concise companion. The README gives the portable build and archival retrieval commands. No underlying response, prompt, reference answer or assessment was deleted.

Codex assisted implementation, annotation, assessment and writing. Qwen generated the experimental facts. Human authors retain responsibility, and these semantic judgments are not independent review. Tool citation: OpenAI (2025), *Introducing Codex*, `https://openai.com/index/introducing-codex/`. The author checklist records unresolved submission-policy questions.