

Inspectable Narrative Graphs: Evidence-Linked Extraction and Contextual Selection

Keywords: Natural Language Processing, Knowledge Representation and Reasoning, Large Language Models, Visualization.

Abstract: We present a narrative extraction workflow connecting readers’ questions, generated relationships, temporal and attribution details, and visual inspection of evidence and errors. We compare extraction before seeing a question followed by fixed selection with independent contextual extraction on three synthetic stories and three short published passages. Assessment separates strict agreement with reference answers from preserved meaning, relevance and unresolved interpretation. All 12 story responses were parseable. Fixed selection matched every possession and dated location target without extras, whereas contextual answers retained irrelevant facts. On fable actions, strict F1 was 0.625 for selection and 0.308 for contextual extraction. Alice’s contextual claims preserved 3 of 4 semantic targets despite zero strict F1. Prose parsing succeeded in 7 of 9 calls. The workflow makes these mixed results inspectable, including useful thought representations, incorrect participants and lost conditional scope. These small development comparisons identify concrete extraction and selection problems for further research. Reference annotations, assessments and manuscript preparation used OpenAI Codex (OpenAI, 2025).

1 INTRODUCTION

We present a workflow for extracting and inspecting narrative relationships in response to a reader’s question. Questions about possession, location or belief call for different views of the same evidence. A useful graph makes relationships accessible while preserving the information needed to check them. Looking into a book is not equivalent to reading it. A character’s belief about an object does not establish its location.

We compare two implemented approaches. *Pre-extract/select* (A) generates a fact collection without seeing the question, then applies fixed filters to its existing records. *Contextual extraction* (B) independently generates facts from the same narrative with the question supplied. Both use readable endpoints, evidence citations, temporal bounds and attribution. Our contributions are:

1. An implemented workflow connecting questions, narratives, generated relationships and qualifications to evidence inspection. Compact records and linked views let readers examine an edge’s assertion and cited text (Section 3 and Figures 1–3).
2. An empirical comparison of extraction before and after question disclosure over identical evidence. Synthetic and prose results help researchers distinguish extraction errors from selection errors (Section 5).

3. Combined assessment and visualization of strict agreement with reference answers, semantic preservation, relevance and unresolved interpretation. Tables 1 and 2 and the annotated figures expose useful partial outputs alongside substantive errors.

Our position is that these dimensions should be assessed together: a single agreement score cannot explain which narrative meanings survive. We study the implemented prototype through small development comparisons.

2 RELATED WORK

Extraction and representation. TextRunner extracts relations without a predefined vocabulary and indexes tuples for queries (Banko et al., 2007), establishing a precedent for extraction before a question. SPIRES recursively prompts a language model to populate a supplied schema and grounds entities in external vocabularies (Caufield et al., 2024). We instead retain local names and evidence spans in a compact format and compare fixed selection with extraction conditioned on a question.

OWL-Time represents instants, intervals, durations and ordering (Cox and Little, 2017). GRaSP separates instances, mentions and attribution, linking propositions to sources and stances (Fokkens et al.,


Orchard · text, model network and fixed query selection

S1 Vera owns Copper Book.


S3 Jun carries Copper Book from day 9 up to but not including day 17.

S11 Jun believes that Ivory Vase is in Hill Court.

Model extraction before the question · 10 nodes / 14 records



Fixed selection: ownership/carrying · 6 nodes / 5 records



Blue fill: person. Cream: object. Grey: place (display categories). loc = located in. Dashed broad edges carry holder / attitude, not extra relationships. + supported. Unprinted bounds and attribution are null. Source excerpted: all 14 sentences supplied.

Figure 1: Orchard: the actual model extraction before the question (top network), followed by fixed selection from that extraction (bottom). Executed question: Who owns or carries what? Return all and only directly narrated ownership and carrying relationships, retaining every stated interval. All 14 extracted records and all five selected records are shown. Source excerpts S1, S3 and S11 come from the complete fourteen-sentence input. Exact endpoint strings share nodes. The selected six-node network is connected. Holder/attitude badges qualify two location edges without adding relationships. The broad extraction is not a gold graph. + denotes retained support for the selected records, not a selection colour. Node categories are display annotations. Horizontal anchors are shared, with compressed vertical spacing below.

2017). Its interface also connects events, sources and text. These precedents motivate our explicit intervals and holder/attitude fields, which address a limited subset of those representational needs.

Comparison and inspection. RelVis integrates extraction comparison, visual inspection of sentences and tuples, quantitative evaluation and error annotation (Schneider et al., 2017). It distinguishes valid

extractions absent from reference annotations from errors. CaRB improves annotations and matching, including partial tuple coverage and checks respecting argument and relation positions (Bhardwaj et al., 2019). Our strict matcher is not CaRB’s scorer. Separate semantic judgments explain preserved meaning, missing qualifications and unresolved readings.

The visual design applies established filtering, linked selection and coordinated views (Heer and Shneiderman, 2012). Our contribution integrates these ideas with narrative questions, temporal and attribution details, and extraction timing relative to the question. Actual outputs expose differences such as a supported broad action with an incorrect participant. Visualization, evidence links and partial credit are not individually new. We neither extend these systems’ implementations nor evaluate them as baselines.

3 REPRESENTATION AND METHODS

3.1 Evidence and Qualified Records

Each request receives the complete selected story or contiguous excerpt with numbered evidence spans and unchanged wording. Output facts contain eight fields: subject, relation, object, evidence_ids, valid_from, valid_until, holder and attitude. The model authors every value. The runtime never fills semantics from reference answers.

For example, the following formatting illustration is authored, not a model result:

```
{ "subject": "Nara", "relation": "carries",
  "object": "map", "evidence_ids": ["X1"],
  "valid_from": 2, "valid_until": 5,
  "holder": null, "attitude": null }
```

Given explicit evidence, this expresses carrying over [2,5), including the start but excluding the end. A query window [3,4) selects this overlapping record without clipping its source bounds. Intrinsic validity concerns when a relationship holds. Occurrence or observation alone does not establish onset or duration. Unsupported bounds remain null: unspecified, not timeless or false. Relative phrases such as “shortly after” are not converted to numbers, so their temporal detail may be lost.

For belief, subject, relation and object express the believed relationship. Holder names the believer and attitude is *believes*. Direct narration uses null holder and attitude without implying character knowledge. A mental or speech act may instead be an explicit relationship, but its uncertain content must not become reality. Pronoun resolution requires a

clear referent in the excerpt. The frozen matcher interprets missing nullable fields as unspecified, while field compliance records their absence.

3.2 Extraction Before or With the Question

Pre-extract/select (A) first requests all explicitly supported relationships without showing the questions. Its actual output is frozen before any contextual request is sent. CPU selection retains existing records by relation category, requested entity, attribution status or explicit interval overlap. Synthetic rules distinguish possession, dated location and belief/narration. Prose rules use lexical action and claim categories appropriate to each passage. These fixed rules do not infer relationships, repair meanings or consult reference answers.

Contextual extraction (B) receives the same complete evidence and one question, requesting only relationships needed to answer it. It never consumes A’s graph. Both approaches share the representation, model and authoring conventions. Prompts, questions, selectors and call order are frozen before each batch. All extractions without questions precede contextual transmission. Each request runs once without adaptive repair, independently of another output’s failure.

The experiments use Qwen/Qwen3-8B-AWQ in nonthinking mode (Yang et al., 2025; Qwen Team, nd), served by vLLM (Kwon et al., 2023) on one RTX 4090. Temperature is 0.7, top-p 0.8, top-k 20, and seed 1988649846. Each request reserves 2,048 output tokens within a context of 12,288 tokens. The pinned revision is recorded in the appendix. Ordinary JSON is requested without a guided schema. Both approaches receive the same short formatting example using unrelated facts, without a tested reference answer.

3.3 Inspecting Actual Output

Graph edges retain generated direction, names, citations and qualifications. Only identical endpoint strings share a display identity. Aliases are not inferred, and objects containing sentences remain literal nodes. The Orchard network makes shared participants and connectivity visible (Figure 1). Node fills denote display categories, not model types or correctness. Status symbols and edge styles distinguish supported, partial, unsupported, unresolved and irrelevant content. A citation highlight is a navigation aid, not verification. Other figures use relationship rows for long labels. Missing targets appear beside graphs,

never as corrected model edges. Display subsets do not change scoring denominators.

4 EVALUATION

4.1 Materials and Annotation

The compact-story v2 comparison uses Harbor, Orchard and Museum, each with 14 evidence sentences. Three questions request possession, location during an interval, and belief contrasted with narration. Harbor and Orchard were used in development. Museum adds different sentence ordering and bindings frozen before this batch. Earlier ladders and historical syntax recoveries are not pooled with these results.

The prose comparison uses Townsend’s complete Lion and Mouse fable (133 words), the opening of *Alice’s Adventures in Wonderland* (112 words), and a contiguous opening passage from *The Red-Headed League* (159 words) (Aesop, 2008; Carroll, 2008; Doyle, 1999). Citations identify the retained Gutenberg editions, not original publication dates. Passages were selected before execution for understandable participants, actions and thought or dialogue. Two questions address actions and claims in each fictional passage.

Evidence, question scope, reference alternatives and normalization were frozen before generation. Reference answers and subsequent assessments of the actual responses were Codex-authored (OpenAI, 2025), not independently validated by humans. Synthetic source inventories distinguish supported but irrelevant records from unsupported ones. Prose reference annotations cover the requested information. Unmatched wording alone is not treated as false. Ambiguous judgments, including Holmes’s vocative and Alice’s daisy-chain consideration, remain open to author review.

4.2 Agreement, Meaning and Failure

A relationship can identify the right participants yet omit when it holds or whose belief it expresses. Our strict score requires the whole record to agree with a reference answer: the normalized relationship, evidence set, temporal bounds and attribution. We call this *strict qualified-fact matching*. Predictions are matched one-to-one to frozen alternatives. Precision is matches divided by all predictions. Recall is matches divided by reference targets, and F1 is their harmonic mean. Incorrect, duplicate, irrelevant and unresolved predictions remain in the denominator. Omissions reduce recall. Bare relationship mea-

sures ignore qualifications and citations. Strict F1 measures agreement with this annotation contract, not the proportion of statements that are factually true.

Semantic assessment separately records supported, partially supported, unsupported and unresolved content, qualification errors and relevance. Target coverage counts distinct target meanings explicitly preserved. One combined object may cover several targets while still being one prediction. Partial support is not full target coverage. Absence of a known contradiction is not evidence of support. These measures expose useful content without inventing missing semantics or changing equivalence rules after seeing which approach benefits.

Strict parsing rejects malformed JSON, duplicate keys and non-JSON constants. A narrowly permitted derived recovery can remove a trailing comma immediately before a closing delimiter outside strings, retaining raw bytes and an edit record. It cannot complete truncated content or remove extra braces. Neither main batch required such recovery. Unparseable results have unavailable graph scores, not the semantics of an intentionally empty graph.

The analysis is descriptive. These are purposively selected development examples, not independent repeated trials of a population effect. Questions sharing a story and outputs reused by selection are dependent. We do not use significance tests or treat individual records as independent experimental units.

5 RESULTS

5.1 Synthetic Selection and Qualifications

All 12 compact-story outputs are strictly parseable. Pre-extract/select (A) matches every possession and dated location target without extras (Table 1). Contextual extraction (B) commonly retrieves the requested physical relationships but returns irrelevant facts in every answer. Orchard’s location answer also omits a target: recall is 0.750. The fresh Museum story shows the same excessive selection as the earlier stories.

Orchard illustrates actual selection from a broader model extraction (Figure 1). Its five ownership/carrying records connect six nodes, retaining explicit carrying intervals. This illustration does not establish a construction effect. The companion contrasts this slice with the twelve-record contextual answer, which also retains seven irrelevant locations.

Among emitted records with identifiable explicitly timed sources, 79 of 79 preserve their full intervals, including irrelevant records. This is conditional

The Lion and the Mouse · actions and rescue


S2 ... "If you would only spare my life, I would be sure to repay your kindness."

S5 The Mouse, recognizing his roar, came and gnawed the rope with his teeth, and set him free, exclaiming:

Request: physical actions, excluding dialogue and intentions

A: text → extraction → selection

B: text + question → extraction



+ supported meaning ~ partial meaning. Citations are not proof of support.

A: spare loses conditional scope. B: early capture/release missing. Source excerpted. Full fable supplied.

Figure 2: Fable actions from pre-extract/select (A) and independent contextual extraction (B). + indicates supported meaning and ~ partial meaning. Display: five of eight A records and all five B records. A's awakening, hunter capture and rope binding are omitted only from display. All predictions remain scored. Executed question: Which physical running/waking, capture, binding, rope-gnawing and release relationships involve the Lion, Mouse, hunters and ropes? Exclude dialogue, intentions, laughter and the moral. Source display excerpts S2 and S5 from the complete supplied fable. The spare edge loses conditional scope and retains the misspelled key `valid until`. B loses the face-specific running endpoint and omits the early capture/release. Other displayed qualifications are null. Assessments are Codex-authored.

preservation, not temporal recall over every reference. The companion's Harbor timeline distinguishes these intervals from the query window. Attribution is less stable: Orchard's pre-extraction decomposes beliefs correctly, whereas Harbor and Museum use sentence-valued belief objects that disrupt selection.

5.2 Published Prose and Partial Meaning

Of 9 prose responses, 7 are strictly parseable. Holmes's pre-extraction and contextual action answer end with an extra closing brace. Their unavailable views remain in Table 2. Failure of the shared pre-extraction affects both A questions. The parseable

Holmes claims answer has partial and unresolved content rather than fully covered targets.

For fable actions, pre-extract/select has strict F1 0.625 versus 0.308 for contextual extraction, and semantic coverage 6 of 8 versus 4 of 8. Both capture the later rescue, but the contextual answer omits the early capture and release. A's unqualified `spare` relation loses the scope of the Mouse's conditional plea (Figure 2). Some supported passive formulations also fail strict matching: semantic preservation and agreement with reference answers are not interchangeable.

Alice shows a different comparison. Contextual claims preserve 3 of 4 semantic targets, compared with 1 of 4 for pre-extract/select, despite zero strict F1 for both. The thought object preserves rhetorical


Alice · inspecting an action and a thought

S1 ... once or twice she had peeped into the book her sister was reading ...

S1 ... "and what is the use of a book," thought Alice "without pictures or conversations?" ...


S2 ... whether the pleasure of making a daisy-chain would be worth the trouble of getting up and picking the daisies ...

B action request: sitting, reading/inspection and running



x The sister reads. Alice peeps. Same source, different participant.

B claims request: absent book content and Alice's thoughts



Final edge: holder = Alice, attitude = null. Other qualifiers null. Thought remains a sentence. Daisy-chain judgment unresolved.

+ supported meaning x unsupported ? unresolved. Source excerpted. Full passage supplied to each independent request.

Figure 3: Alice: independent contextual extraction (B) for two requests over identical complete evidence. Display: two of five action edges (two sitting edges and Rabbit running omitted) and two of three claims (book content omitted). All remain scored. Executed action question: Which sitting/location, reading/inspection and running relationships involve Alice, her sister, the book, bank and Rabbit? Exclude thoughts, feelings, appearance and book-content properties. Executed claims question: What does the passage say about the book's absent content and Alice's thoughts about the usefulness of such books and making a daisy-chain? Exclude physical actions, locations, appearance and the Rabbit; preserve questions and consideration without treating them as decisions. Source excerpts are marked. + indicates supported meaning, x an unsupported participant, and ? unresolved consideration. The thought remains the actual sentence-valued object. All bounds are null. The final record has holder Alice and null attitude.

meaning but not the requested decomposition. Contextual extraction also assigns reading to Alice instead of her sister. The daisy-chain record leaves open whether considering worth has become judging it worthwhile (Figure 3). More covered meaning can therefore accompany a substantive error or unresolved interpretation.

5.3 Cost of the Presented Experiments

The two main batches comprise 21 generations and 454.82 allocated GPU seconds. Compact stories used 249.45 seconds and prose 205.37 seconds. Summed request times were 104.72 and 57.24 seconds respectively. Allocation includes loading, checks, service idle time and shutdown. Reused pre-extractions count once. These batch costs exclude earlier development and do not measure production latency.

Table 1: Compact-story v2 strict qualified-fact scores. A: pre-extract/select. B: contextual extraction. n includes every prediction. Targets per story: possession 5, location 4, belief/reality 4. No historical recoveries are pooled.

Story	Question	n_A	P_A	R_A	$F1_A$	n_B	P_B	R_B	$F1_B$
Harbor	Possession	5	1.000	1.000	1.000	10	0.500	1.000	0.667
Harbor	Location	4	1.000	1.000	1.000	6	0.667	1.000	0.800
Harbor	Belief/reality	5	0.000	0.000	0.000	14	0.143	0.500	0.222
Orchard	Possession	5	1.000	1.000	1.000	12	0.417	1.000	0.588
Orchard	Location	4	1.000	1.000	1.000	5	0.600	0.750	0.667
Orchard	Belief/reality	4	1.000	1.000	1.000	13	0.154	0.500	0.235
Museum	Possession	5	1.000	1.000	1.000	13	0.385	1.000	0.556
Museum	Location	4	1.000	1.000	1.000	5	0.800	1.000	0.889
Museum	Belief/reality	4	0.000	0.000	0.000	14	0.143	0.500	0.222

Table 2: Published prose: strict agreement and Codex-authored semantic assessment. A: pre-extract/select. B: contextual extraction. S/P/X/U counts supported, partial, unsupported and unresolved records. Coverage counts preserved target meanings. NA denotes parse failure, not an empty graph. All parseable predictions cite valid evidence IDs, which does not establish support.

Passage	Question	Method	n	Precision	Recall	F1	S/P/X/U	Coverage
Alice	Actions	A	2	0.500	0.200	0.286	2/0/0/0	3/5
Alice	Actions	B	5	0.400	0.400	0.400	3/1/1/0	3/5
Alice	Claims	A	4	0.000	0.000	0.000	2/2/0/0	1/4
Alice	Claims	B	3	0.000	0.000	0.000	2/0/0/1	3/4
Fable	Actions	A	8	0.625	0.625	0.625	6/1/0/1	6/8
Fable	Actions	B	5	0.400	0.250	0.308	4/1/0/0	4/8
Fable	Claims	A	1	0.000	0.000	0.000	1/0/0/0	2/5
Fable	Claims	B	4	0.000	0.000	0.000	1/3/0/0	0/5
Holmes	Actions	A	NA	NA	NA	NA	NA	NA
Holmes	Actions	B	NA	NA	NA	NA	NA	NA
Holmes	Claims	A	NA	NA	NA	NA	NA	NA
Holmes	Claims	B	3	0.000	0.000	0.000	0/1/0/2	0/5

6 DISCUSSION AND LIMITATIONS

The workflow enables inspection of both useful information and errors. The fable exposes a conditional plea rendered as an unqualified relationship. Alice juxtaposes preserved thought with an incorrect reading participant. Harbor shows intact temporal bounds alongside irrelevant facts. These examples connect measured discrepancies to the source wording and the requested information, giving researchers concrete targets for improving extraction and selection.

Fixed selection cannot recover an omitted fact or repair attribution. In the fable claims view, its exact subject filter expects “Lion” and rejects “A LION”, discarding useful records. In Alice, feeling records embedding actions can be excluded by the action filter. The results therefore concern these particular implementations, including brittle selectors. Conversely, repeated irrelevant facts in contextual answers show that a prominent question and compact format alone do not ensure relevance.

The small, purposive samples limit generalization. Several examples were reused during development, familiar prose may have been in model training, and there is one generation per request. Shared AI assistance in implementation, annotation and assessment creates possible correlated blind spots. Reference coverage and semantic interpretations require further review. The compact fields cannot faithfully encode every relative time phrase, nested perspective or event structure. No participant evaluation establishes usability, and differences between these extracted views do not demonstrate query-dependent ontology construction.

The next research steps are independent review of reference coverage, evaluation on fresh frozen passages, and targeted improvements in participant, attribution and relevance handling without supplying passage-specific answers. Reader studies can then test whether the inspection workflow improves comprehension and error detection.

7 CONCLUSION

We present a workflow connecting narrative questions, qualified relationships and evidence inspection. Comparing pre-extract/select with contextual extraction reveals useful action and thought information alongside relevance and participant errors. Selection performs better on several questions, while contextual extraction preserves more of Alice’s claims. Combining strict agreement, semantic assessment and visualization provides a concrete basis for improving narrative extraction and studying its use.

8 AI USE AND ANNOTATION PROVENANCE

OpenAI Codex (OpenAI, 2025) assisted implementation, analysis, rendering code, reference annotation, semantic assessment and writing all sections, including the abstract and captions. Figures deterministically display retained records. Qwen generated the experimental facts. The Codex tool citation does not identify its historical backend. Its judgments are not independent human validation. Pending decisions remain unresolved, and human authors retain responsibility.

REFERENCES

Aesop (2008). *Three Hundred Aesop’s Fables*. Project Gutenberg, ebook 21. G. F. Townsend, trans. Digital edition released 25 June 2008, updated 15 October 2025. Retrieved 9 September 2026. <https://www.gutenberg.org/ebooks/21>.

Banko, M., Cafarella, M. J., Soderland, S., Broadhead, M., and Etzioni, O. (2007). Open information extraction from the web. In *Proceedings of the 20th International Joint Conference on Artificial Intelligence (IJCAI)*, pages 2670–2676.

Bhardwaj, S., Aggarwal, S., and Mausam (2019). CaRB: A crowdsourced benchmark for open IE. In *Proceedings of EMNLP-IJCNLP*, pages 6262–6267. Association for Computational Linguistics. <https://doi.org/10.18653/v1/D19-1651>.

Carroll, L. (2008). *Alice’s Adventures in Wonderland*. Project Gutenberg, ebook 11. Digital edition released 27 June 2008, updated 26 June 2025. Chapter I. Retrieved 9 September 2026. <https://www.gutenberg.org/ebooks/11>.

Caufield, J. H., Hegde, H., Emonet, V., Harris, N. L., Joachimiak, M. P., Matentzoglou, N., Kim, H., Moxon, S., Reese, J. T., Haendel, M. A., Robinson, P. N., and Mungall, C. J. (2024). Structured prompt interrogation and recursive extraction of semantics (SPIRES):

a method for populating knowledge bases using zero-shot learning. *Bioinformatics*, 40(3):btae104. <https://doi.org/10.1093/bioinformatics/btae104>.

Cox, S. and Little, C., editors (2017). *Time Ontology in OWL*. World Wide Web Consortium. W3C Recommendation, 19 October 2017. <https://www.w3.org/TR/2017/REC-owl-time-20171019/>.

Doyle, A. C. (1999). *The Adventures of Sherlock Holmes*. Project Gutenberg, ebook 1661. Digital edition released 1 March 1999, updated 10 October 2023. The Red-Headed League. Retrieved 9 September 2026. <https://www.gutenberg.org/ebooks/1661>.

Fokkens, A., Vossen, P., Rospocher, M., Hoekstra, R., and van Hage, W. R. (2017). GRASP: Grounded representation and source perspective. In *Proceedings of the Workshop Knowledge Resources for the Socio-Economic Sciences and Humanities associated with RANLP 2017*, pages 19–25. INCOMA Inc. https://doi.org/10.26615/978-954-452-040-3_003.

Heer, J. and Shneiderman, B. (2012). Interactive dynamics for visual analysis. *Communications of the ACM*, 55(4):45–54. <https://doi.org/10.1145/2133416.2146416>.

Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., Gonzalez, J. E., Zhang, H., and Stoica, I. (2023). Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*, pages 611–626. Association for Computing Machinery. <https://doi.org/10.1145/3600006.3613165>.

OpenAI (2025). Introducing Codex. <https://openai.com/index/introducing-codex/>. 16 May. Tool announcement.

Qwen Team (n.d.). Qwen3-8B-AWQ model card and pinned snapshot. Hugging Face, pinned snapshot. Revision specified in appendix. Accessed 9 September 2026.

Schneider, R., Oberhauser, T., Klatt, T., Gers, F. A., and Löser, A. (2017). Analysing errors of open information extraction systems. In *Proceedings of the First Workshop on Building Linguistically Generalizable NLP Systems*, pages 11–18. Association for Computational Linguistics. <https://doi.org/10.18653/v1/W17-5402>.

Yang, A., Li, A., Yang, B., et al. (2025). Qwen3 technical report. arXiv:2505.09388.

APPENDIX

Qwen3-8B-AWQ revision:
4da05a8edb55c6046cce958586c33b61da07bb79.
Software: vLLM 0.10.2, Transformers 4.55.2. The pinned template uses `enable_thinking=False`. The concise companion indexes complete machine-readable requests, outputs, reference alternatives and assessments.