

Supplementary material: retained exploratory narrative graphs

This supplement preserves the eight original review plates and the full-output appendix without altering their PDF pages. The new main figures use labelled display subsets; all evaluated records and historical failures remain available. References and semantic assessments are Codex-authored, not independent human review.

Contents and evaluation versions

The pages below provide detailed canonical tables and exact executed questions. The original eight-plate review document follows, then its full-output appendix, including unparseable Holmes outputs shown as raw text. Source passages, reference alternatives, and all calls are also available in the immutable repository reports linked in the manuscript manifest.

S1. Compact stories v2: fresh outputs only

Story	Question	A n	A P/R/F1	B n	B P/R/F1
Harbor	possession	5	1.000/1.000/1.000	10	0.500/1.000/0.667
Harbor	locations	4	1.000/1.000/1.000	6	0.667/1.000/0.800
Harbor	belief	5	0.000/0.000/0.000	14	0.143/0.500/0.222
Orchard	possession	5	1.000/1.000/1.000	12	0.417/1.000/0.588
Orchard	locations	4	1.000/1.000/1.000	5	0.600/0.750/0.667
Orchard	belief	4	1.000/1.000/1.000	13	0.154/0.500/0.235
Museum	possession	5	1.000/1.000/1.000	13	0.385/1.000/0.556
Museum	locations	4	1.000/1.000/1.000	5	0.800/1.000/0.889
Museum	belief	4	0.000/0.000/0.000	14	0.143/0.500/0.222

S2. Published prose: separate strict and semantic outcomes

S/P/X/U denotes supported, partial, unsupported and unresolved records. Coverage is unique explicit target concepts. NA is unavailable parsing, not an empty graph.

Passage / question	Method	n	Strict P/R/F1	S/P/X/U	Coverage
Alice / actions	A	2	0.500/0.200/0.286	2/0/0/0	3/5
Alice / actions	B	5	0.400/0.400/0.400	3/1/1/0	3/5
Alice / claims	A	4	0.000/0.000/0.000	2/2/0/0	1/4
Alice / claims	B	3	0.000/0.000/0.000	2/0/0/1	3/4
Fable / actions	A	8	0.625/0.625/0.625	6/1/0/1	6/8
Fable / actions	B	5	0.400/0.250/0.308	4/1/0/0	4/8
Fable / claims	A	1	0.000/0.000/0.000	1/0/0/0	2/5
Fable / claims	B	4	0.000/0.000/0.000	1/3/0/0	0/5
Holmes / actions	A	NA	NA/NA/NA	NA	NA
Holmes / actions	B	NA	NA/NA/NA	NA	NA
Holmes / claims	A	NA	NA/NA/NA	NA	NA
Holmes / claims	B	3	0.000/0.000/0.000	0/1/0/2	0/5

S3. Distinct batch costs

Service allocation includes loading, checks, inference and shutdown. Request times alone exclude those costs. Costs may be added; accuracy across incompatible versions may not.

Batch (not pooled)	Calls	Input / output tokens	Request s	Allocated s
Simple ladder	8	1344 / 848	20.03	157.35
Simple ladder v2	8	2594 / 997	22.45	171.37
Compact stories v2	12	8088 / 5860	104.72	249.45
Published prose	9	6802 / 2927	57.24	205.37

The two main comparison batches used 454.82 allocated seconds. Historical whole-project allocation at their conclusion was 10717.91 seconds, including earlier attempts and failures. No new inference was performed for this manuscript. Small-batch timing does not establish production-tail latency or feasibility of the larger registered study.

S4. Development context

The original simple ladder recovered its tested direct facts, contextual selections and explicit intervals before encountering belief decomposition and citation/interval failures. The revised ladder retained its regression successes and handled original and fresh simple beliefs; office examples remained imperfect. The office questions permit the same office-mediated organization and do not establish an ontology-construction contrast. Earlier unsuccessful comprehensive-interface attempts motivated simplification but are not compact-interface results. Scores below remain separate by version.

ladder — unchanged development scores

Case	Task	Precision	Recall	F1
1	direct	1.000	1.000	1.000
2	direct	1.000	1.000	1.000
3	selection	1.000	1.000	1.000
4	selection	1.000	1.000	1.000
5	temporal	1.000	1.000	1.000
6	epistemic	0.000	0.000	0.000
7	ontology_contrast	0.000	0.000	0.000
8	ontology_contrast	0.500	0.500	0.500

ladder v2 — unchanged development scores

Case	Task	Precision	Recall	F1
1	direct	1.000	1.000	1.000
2	temporal	1.000	1.000	1.000
3	epistemic	1.000	1.000	1.000
4	epistemic	1.000	1.000	1.000
5	office-person	0.000	0.000	0.000
6	office-continuity	0.000	0.000	0.000
7	office-person	0.500	0.500	0.500
8	office-continuity	0.500	0.500	0.500

S5. Historical syntax recovery is not new generation

Earlier Harbor/Orchard possession outputs failed strict parsing. Their separately retained comma-only derivations are below. Original strict failures are unchanged and excluded from the fresh table. Full edit records and original scores are reproduced unchanged in tables/historical_syntax_recovery.json.

Story	Matches	Predictions	References	P/R/F1
harbor	3	10	5	0.300/0.600/0.400
orchard	3	10	5	0.300/0.600/0.400

S6. Exact requests, display notes and source attribution

Question headlines in the artwork abbreviate the executed requests below. A means fixed selection from an actual query-blind extraction; B is independently generated from text plus question. These are comparisons, not a transformation from A into B.

Figure 1: complete question and display notes

Figure 1. Aesop's The Lion and the Mouse (Townsend translation). Exact executed question: "Which physical running/waking, capture, binding, rope-gnawing and release relationships involve the Lion, Mouse, hunters and ropes? Exclude dialogue, intentions, laughter and the moral." A selects from an actual pre-extraction; B is generated independently. All B action records and the A capture, conditional-spare, release, hunter-capture, gnawing and rescue records are shown. A's waking and rope-binding records are omitted only from display (two of eight); full graphs and all scores remain in the supplement. B preserves the later rescue but omits the earlier capture and release; its running relation loses the face-specific endpoint. A's spare claim loses conditional scope. S6, the final speech, is not displayed but was supplied. Faceted arrows reproduce exact generated endpoint strings; repeated names denote the same exact-string node. Undisplayed qualification values are null except the explicitly noted missing key. Colours/symbols reproduce Codex-authored assessments, not independent verification.

Figure 2: complete question and display notes

Figure 2. Harbor location comparison. Exact executed question: "Which directly narrated location relationships have an explicitly stated interval overlapping day 2 up to but not including day 4? Return all of them with their full evidence-stated intervals, not clipped to the question window. Exclude undated locations and beliefs." All four A location records and the four corresponding B records are displayed. B's other two records are reproduced in the labelled carrying callout and remain in its six-prediction denominator. Both methods preserve source intervals; the shaded question window selects overlap and does not redefine validity. Circles distinguish included starts from excluded ends. Only S5-S8 are printed; the model received the same complete fourteen-sentence story in every call. A is fixed selection; B is independent generation. All generated holder/attitude values in these records are null. This is an extraction/relevance comparison, not ontology construction.

Figure 3: complete question and display notes

Figure 3. Carroll's Alice opening passage, supplied unchanged to both B requests. Exact action question: "Which sitting/location, reading/inspection and running relationships involve Alice, her sister, the book, bank and Rabbit? Exclude thoughts, feelings, appearance and book-content properties." Exact claims question: "What does the passage say about the book's absent content and Alice's thoughts about the usefulness of such books and making a daisy-chain? Exclude physical actions, locations, appearance and the Rabbit; preserve questions and consideration without treating them as decisions." Three of five action records show inspection, incorrect assignment of reading to Alice, and the Rabbit's literal object 'close by her'; the two sitting records remain in the full graph and score. All three claims records are shown. The sentence-valued thinks object preserves rhetorical meaning but not the requested decomposed attribution structure; it is not redrawn as a corrected proposition. The daisy-chain record lacks attitude alongside holder Alice and may turn consideration into a judgment. All generated bounds are null; other holder/attitude fields are null. Display line breaks and underscore-to-space predicate typography do not change identities, meanings or evaluations. Status labels reproduce the retained Codex assessment.

Alice's Adventures in Wonderland

Lewis Carroll. Chapter I: Down the Rabbit-Hole, opening two paragraphs. [Source](#). Retrieval: 2026-09-09T13:42:39.616173+00:00. Excerpt SHA-256: 63456811be1ee205a9908656ee756338064caa23145f1f8829dcb0ec52f1f77c.

Project Gutenberg lists this edition as public domain in the USA; original attribution and complete supplied license notices retained. Check local law outside the USA. Notices: [data/published_prose/pg11_notices.txt](#).

S1 Alice was beginning to get very tired of sitting by her sister on the bank, and of having nothing to do: once or twice she had peeped into the book her sister was reading, but it had no pictures or conversations in it, “and what is the use of a book,” thought Alice “without pictures or conversations?”

S2 So she was considering in her own mind (as well as she could, for the hot day made her feel very sleepy and stupid), whether the pleasure of making a daisy-chain would be worth the trouble of getting up and picking the daisies, when suddenly a White Rabbit with pink eyes ran close by her.

Three hundred Aesop’s fables

Aesop; translated by George Fyler Townsend. The Lion And The Mouse. [Source](#). Retrieval: 2026-09-09T13:42:39.348162+00:00. Excerpt SHA-256: 340385bf6743191af33313e4998e9e16987d7f92d658417861516ce220b25fb0.

Project Gutenberg lists this edition as public domain in the USA; original attribution and complete supplied license notices retained. Check local law outside the USA. Notices: [data/published_prose/pg21_notices.txt](#).

S1 A LION was awakened from sleep by a Mouse running over his face.

S2 Rising up angrily, he caught him and was about to kill him, when the Mouse piteously entreated, saying: “If you would only spare my life, I would be sure to repay your kindness.”

S3 The Lion laughed and let him go.

S4 It happened shortly after this that the Lion was caught by some hunters, who bound him by strong ropes to the ground.

S5 The Mouse, recognizing his roar, came and gnawed the rope with his teeth, and set him free, exclaiming:

S6 “You ridiculed the idea of my ever being able to help you, not expecting to receive from me any repayment of your favour; now you know that it is possible for even a Mouse to confer benefits on a Lion.”

The Adventures of Sherlock Holmes

Arthur Conan Doyle. II. The Red-Headed League, opening visit through the gentleman’s greeting. [Source](#). Retrieval: 2026-09-09T13:42:39.964188+00:00. Excerpt SHA-256: 07b8b4e4daa61ba653e966da4f78816bd72f92eabdfc113ddffb56c33829e024.

Project Gutenberg lists this edition as public domain in the USA; original attribution and complete supplied license notices retained. Check local law outside the USA. Notices: [data/published_prose/pg1661_notices.txt](#).

S1 I had called upon my friend, Mr. Sherlock Holmes, one day in the autumn of last year and found him in deep conversation with a very stout, florid-faced, elderly gentleman with fiery red hair.

S2 With an apology for my intrusion, I was about to withdraw when Holmes pulled me abruptly into the room and closed the door behind me.

S3 “You could not possibly have come at a better time, my dear Watson,” he said cordially.

S4 “I was afraid that you were engaged.”

S5 “So I am. Very much so.”

S6 “Then I can wait in the next room.”

S7 “Not at all. This gentleman, Mr. Wilson, has been my partner and helper in many of my most successful cases, and I have no doubt that he will be of the utmost use to me in yours also.”

S8 The stout gentleman half rose from his chair and gave a bob of greeting, with a quick little questioning glance from his small fat-encircled eyes.

S7. Reproducibility and author review

The pinned model is Qwen/Qwen3-8B-AWQ, revision 4da05a8edb55c6046cce958586c33b61da07bb79, seed 1988649846. Full model/template metadata and call settings originate in the frozen manifests. All outputs, timestamps, request hashes and assessments are linked by manuscript_manifest.json. Figure edge manifests preserve original record dictionaries and statuses. Reference details were checked against primary sources on 9 September 2026.

[AUTHOR_REVIEW.md](#) provides the prioritized evidence-linked judgment questions and the single list of unresolved publication details. It is a request for author decisions, not a completed human review. Any subsequent annotation or score change must be separately versioned. [README.md](#) gives the portable regeneration command. The original visual-review and full-output pages follow unchanged.